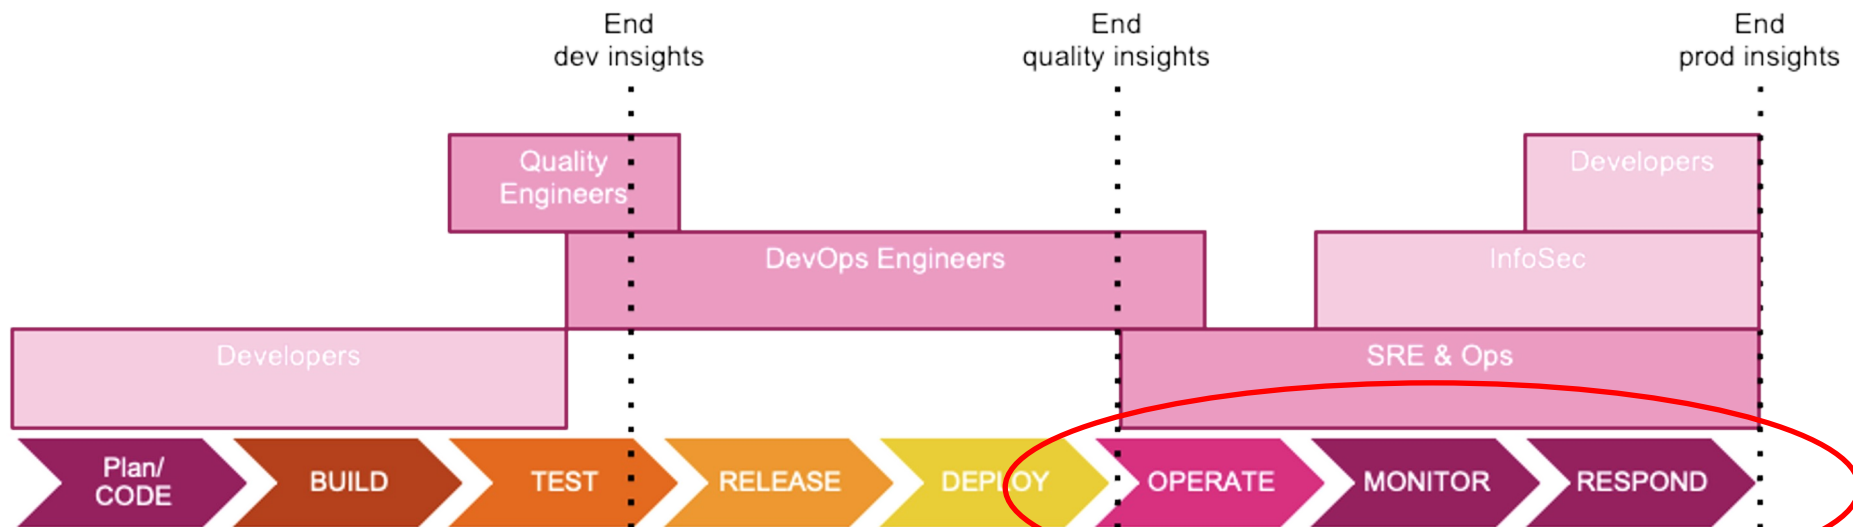


Getting started with Observability & OpenTelemetry

Mar 20, 2022

Who and When in Agile Ops



Area of focus today!

Agenda

- Observability?
- OpenTelemetry?
- Demo in Action!

Introduction



 wwongpairoj@splunk.com

 [wwongpai](#)

 [donler](#)


Warach Wongpairoj

Senior Sales Engineer, Splunk

Previously:

- Senior Sales Engineer, AppDynamics
- Sales Engineer, Oracle Cloud (Cloud Manageability & Cloud Security)

Who and When in Agile Ops



Area of focus today!

“ *Hope is not a strategy.* ”

Traditional SRE saying

It is a truth universally acknowledged that systems do not run themselves. How, then, should a system—particularly a complex computing system that operates at a large scale—be run?

Goal & Responsibility

To address how a complex computing system operating at a large scale should run to achieve scalable and highly reliable software systems.

Responsible for the availability, latency, performance, efficiency, change management, monitoring, emergency response, and capacity planning of their service(s).

Goal & Responsibilityin reality

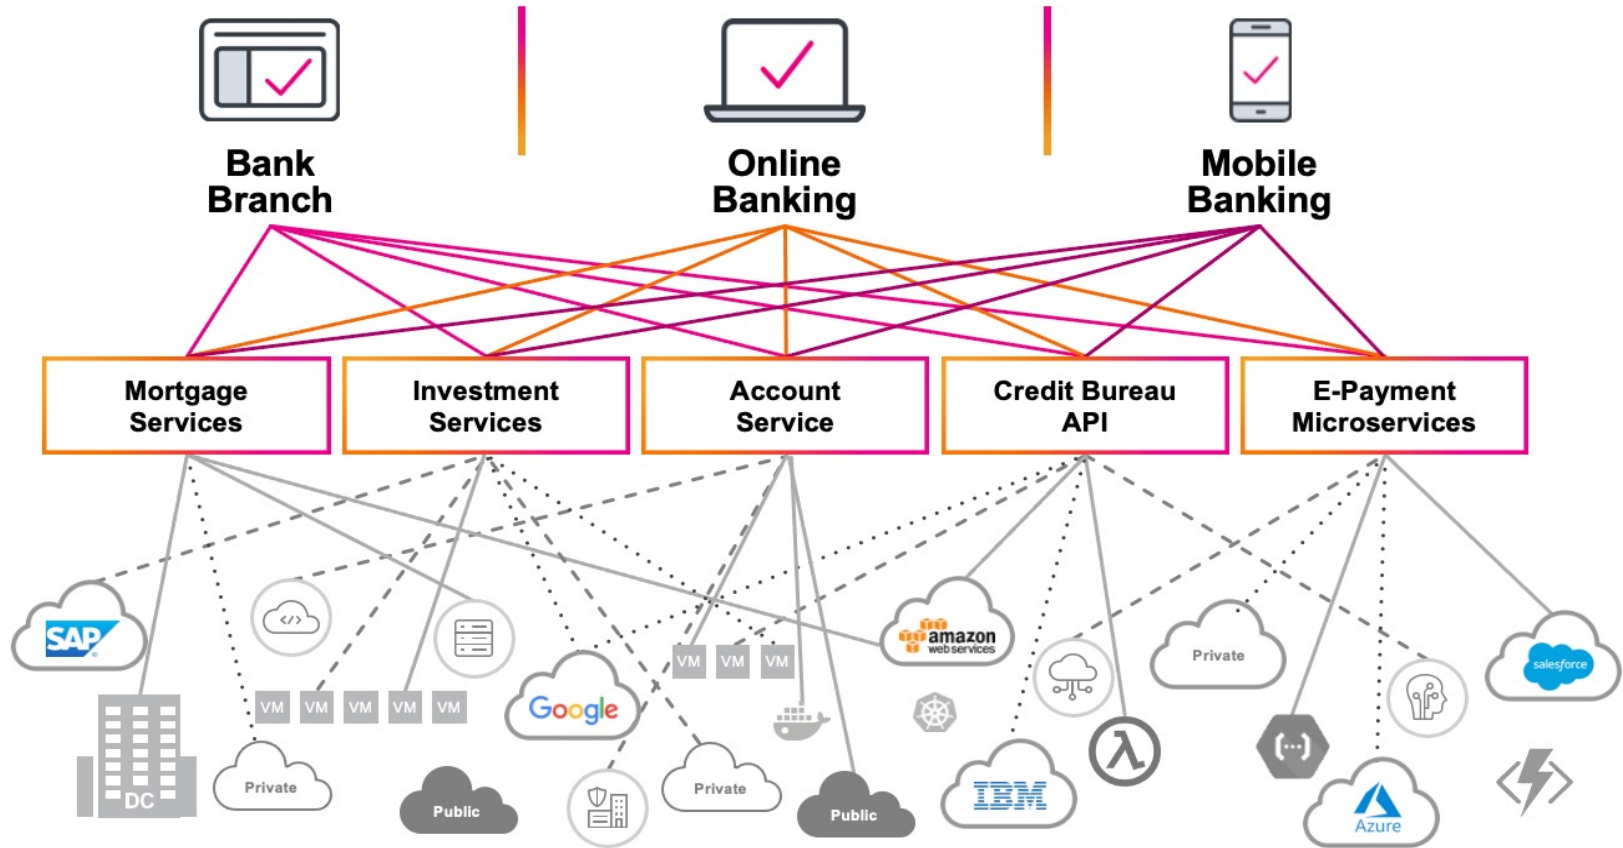
To know who to **wake up** and how to **find root cause**

Which **customer** or action is **impacted**?

MTTD, MTTI, MTTR, Error Budget
and so on.....



Challenges



To Summarize Challenges --> Why Observability is needed

- Microservices create complex interactions
- To manage system at scale is difficult
- Failures don't exactly repeat
- Debugging multi-tenancy is painful
- Traditional monitoring approach can't save us



What is Observability?

- We need to answer questions about our systems.

What characteristics did the queries that timed out at 500ms share in common? Service versions? Browser plugins?

- Instrumentation produces data.
- Querying data answers our questions.

What difference btw Monitoring & Observability?

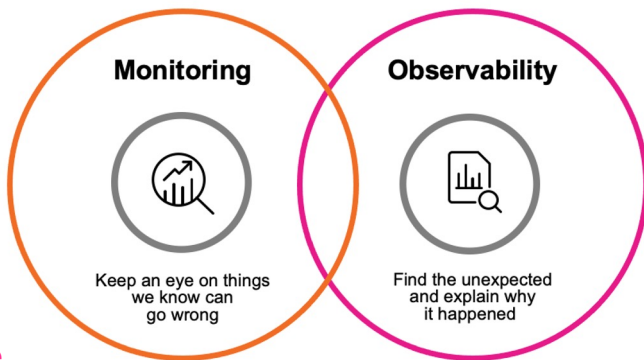
Is it just a fancy word for “Monitoring”?

Looking for expected

problems, e.g.:

- Overloaded CPU
- High memory utilization
- Disk space
- High response latency
- High error rate
- Service availability

Alerting on issues that have occurred.



Looking for unknown unknowns,

e.g.:

- Why are the alerts firing?
- What is in common between problem areas?
- Find the “needle in the haystack”
- Troubleshooting solution for complex systems

Diagnosing why issues have occurred.

But how do I implement these?

- You need data (3 pillars of observability data)
- You need an instrumentation framework!
- and a place to send the data!
- and a way to visualize the data!

LOGS

Detailed debugging information emitted by processes

Data format: LOGS

OSS: Graylog, Elastic,
Fluentd, Fluentbit

- A time stamped text record
- Can be structured (recommended) or unstructured
- Enhanced with metadata

```
{
  "currency": "JPY",
  "http.req.id": "ba09b8be-555f-4da8-9ca3-1519179ab800",
  "http.req.method": "GET",
  "http.req.path": "/product/6E92ZMYFZ",
  "id": "6E92ZMYFZ",
  "message": "serving product page",
  "session": "8867691b-0aa9-4c9c-bfd5-9936859a6db7",
  "severity": "debug",
  "span_id": "6faa8ab4504e5e2a",
  "timestamp": "2020-10-19T20:16:19.324573978Z",
  "trace_id": "6faa8ab4504e5e2a"
}

{
  "http.req.id": "ba09b8be-555f-4da8-9ca3-1519179ab800",
  "http.req.method": "GET",
  "http.req.path": "/product/6E92ZMYFZ",
  "http.resp.bytes": 8726,
  "http.resp.status": 200,
  "http.resp.took_ms": 5,
  "message": "request complete",
  "session": "8867691b-0aa9-4c9c-bfd5-9936859a6db7",
  "severity": "debug",
  "timestamp": "2020-10-19T20:16:19.33012145Z"
}

{
  "http.req.id": "65c08731-c7b7-492b-9a5a-c379509c6842",
  "http.req.method": "GET",
  "http.req.path": "/product/ZZYFJ3GM2N",
  "message": "request started",
  "session": "a80d1d7f-8c5a-4ca3-baaf-a295109f730b",
  "severity": "debug",
  "timestamp": "2020-10-19T20:16:19.620079929Z"
}

{
  "currency": "JPY",
  "http.req.id": "65c08731-c7b7-492b-9a5a-c379509c6842",
  "http.req.method": "GET",
  "http.req.path": "/product/ZZYFJ3GM2N",
  "id": "ZZYFJ3GM2N",
  "message": "serving product page",
  "session": "a80d1d7f-8c5a-4ca3-baaf-a295109f730b",
  "severity": "debug",
  "span_id": "2105a09f662c2df7",
  "timestamp": "2020-10-19T20:16:19.620159918Z",
  "trace_id": "2105a09f662c2df7"
}

{
  "http.req.id": "65c08731-c7b7-492b-9a5a-c379509c6842",
  "http.req.method": "GET",
  "http.req.path": "/product/ZZYFJ3GM2N",
  "http.resp.bytes": 8708,
  "http.resp.status": 200,
  "http.resp.took_ms": 5,
  "message": "request complete",
  "session": "a80d1d7f-8c5a-4ca3-baaf-a295109f730b",
  "severity": "debug",
  "timestamp": "2020-10-19T20:16:19.625605455Z"
}

{
  "http.req.id": "65b65c65-5990-40fb-af6e-a8be4feba183",
  "http.req.method": "GET",
  "http.req.path": "/product/0PUK6V6EV0",
  "message": "request started",
  "session": "8b917480-333b-4e4e-9c39-aeba2bca8c5c",
  "severity": "debug",
  "timestamp": "2020-10-19T20:16:19.82898198Z"
}

{
  "currency": "USD",
  "http.req.id": "65b65c65-5990-40fb-af6e-a8be4feba183",
  "http.req.method": "GET",
  "http.req.path": "/product/0PUK6V6EV0",
  "id": "0PUK6V6EV0",
  "message": "serving product page",
  "session": "8b917480-333b-4e4e-9c39-aeba2bca8c5c",
  "severity": "debug",
  "span_id": "75744bd42bafb568",
  "timestamp": "2020-10-19T20:16:19.82907272Z",
  "trace_id": "75744bd42bafb568"
}

{
  "http.req.id": "65b65c65-5990-40fb-af6e-a8be4feba183",
  "http.req.method": "GET",
  "http.req.path": "/product/0PUK6V6EV0",
  "http.resp.bytes": 8651,
  "http.resp.status": 200,
  "http.resp.took_ms": 5,
  "message": "request complete",
  "session": "8b917480-333b-4e4e-9c39-aeba2bca8c5c",
  "severity": "debug",
  "timestamp": "2020-10-19T20:16:19.834710332Z"
}

2020-10-19 20:16:20 +0000 [info]: #0 stats - namespace_cache_size: 2, pod_cache_size: 19, namespace_cache_api_updates: 878, pod_cache_api_updates: 936, id_cache_miss: 878, pod_cache_host_updates: 19, pod_watch_gone_errors: 1, pod_watch_gone_notices: 1

{"severity": "info", "time": 1603138580, "pid": 1, "hostname": "currency-service-5b468487bd-7q4rj", "name": "currency-service-server", "message": "Getting supported currencies...", "v": 1}
```

The Golden Signals/Telemetry

Google's Golden Signals

Latency, Saturation, Errors, Traffic

USE Monitoring

Utilization, Saturation, Errors

RED Monitoring

Rate, Errors, Duration

USE

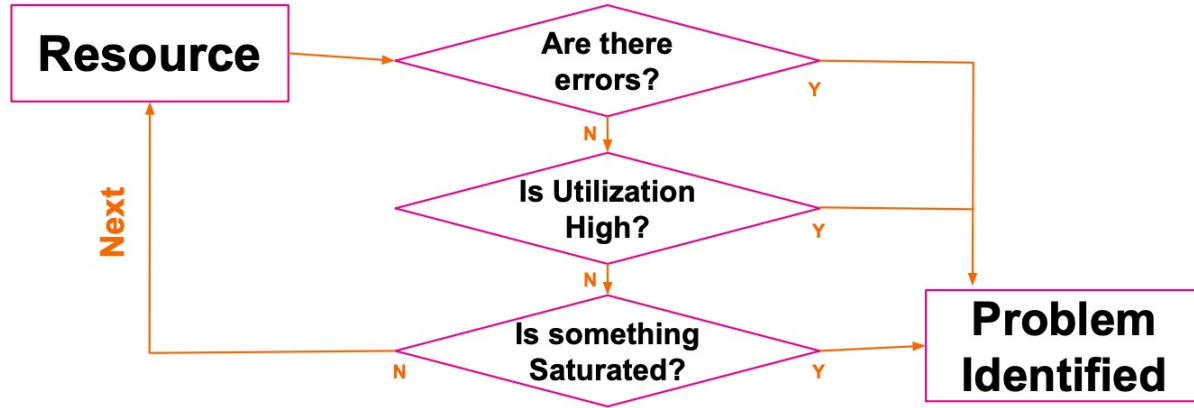
Utilization, Saturation, Error

For every resource, check:

- Utilization
How busy is the resource or amount in use
- Saturation
How much extra work is not being process due to lack of resource
- Errors
All errors

USE

Utilization, Saturation, Error



Data format: **METRICS**
(Time-Series Metrics)

OSS: Prometheus,
Zabbix, Nagios, Cacti,
Collectd etc.

So How About RED?

Rate, Error, Duration

- Designed for request-driven systems, microservices

Rate

- Rate: number/size of requests on network and system
 - HTTP, SOAP, REST
 - Middleware messaging/queuing
 - API calls
 - Overhead of control structures
- Any environment that can fail on peak traffic is a target for rate monitoring

Errors

- Errors: problems that cause an incorrect, incomplete or unexpected result
 - Code failures
 - Production load bugs
 - Peak load bugs
 - Communication woes
- Errors need:
 - Rapid Responses
 - Point Specific responses
- Need deep dive, high-fidelity

Duration

- **Bring events into causal order**
- Both client-side and server-sides are important
 - But client side maybe more
 - Usually (now) the domain of distributed request tracing, RUM and APM

Rate, Error, Duration



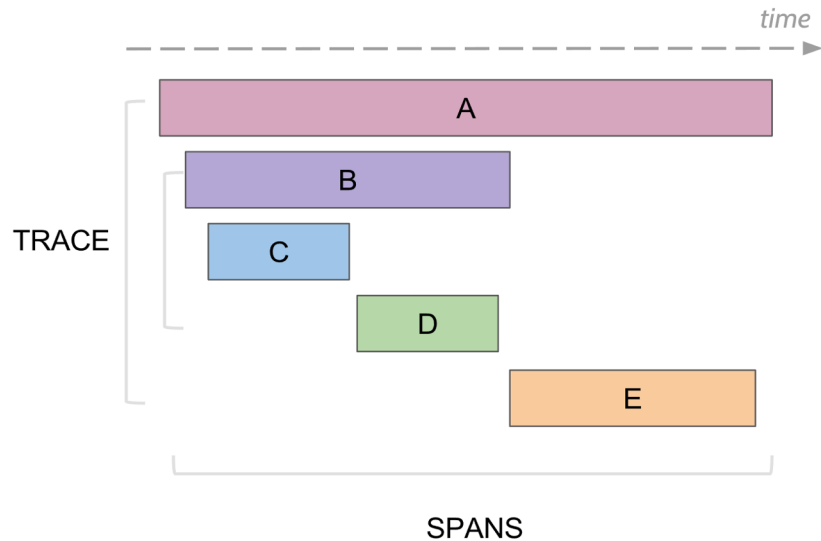
OSS: OpenTracing, Jaeger, Zipkin etc.

RED == Distributed Tracing


Traces

Distributed tracing traverses process, network and security boundaries

- Tracing is a way to record all the operations involved in the handling of a request or transaction through the entire application stack and backend infrastructure
- Spans record individual operations or RPCs, in particular their name, how long they took and where they took place
- Distributed Context contains tracing identifiers, tags, and options that propagated from parent to child spans



Finally, Log, Metric (USE), Trace (RED)..... Done?



Challenges

- Lack of standardization
- Some are vendor-lockin
- Lack of data portability
- The burden on the user to maintain the instrumentation
- Difficult to fix issue with 3 format of data are isolated
- No correlation & Causation


- You need an instrumentation framework!
- Consolidate data (from different format and correlate all logs, metrics and traces)
- Open-standard semantic conventions
- Context-propagation
- Vendor-agnostic



What's OpenTelemetry

The Cloud Native Computing Foundation (CNCF) [OpenTelemetry](#) project provides a single set of open source APIs, libraries, and agents to collect and correlate distributed traces, metrics and logs.

With OpenTelemetry you can instrument your application in a vendor-agnostic way, and then analyze the telemetry data in your backend tool of choice, whether Splunk ,Prometheus, Jaeger or others. Instrument once, and use it anywhere.





**OpenTelemetry
is the second
most active
project in CNCF
today! 🎉**

(per CNCF DevStats)

Everyone is Contributing and Adopting

Cloud Providers



AWS | Azure | GCP

Vendors




Every major vendor!

End-users



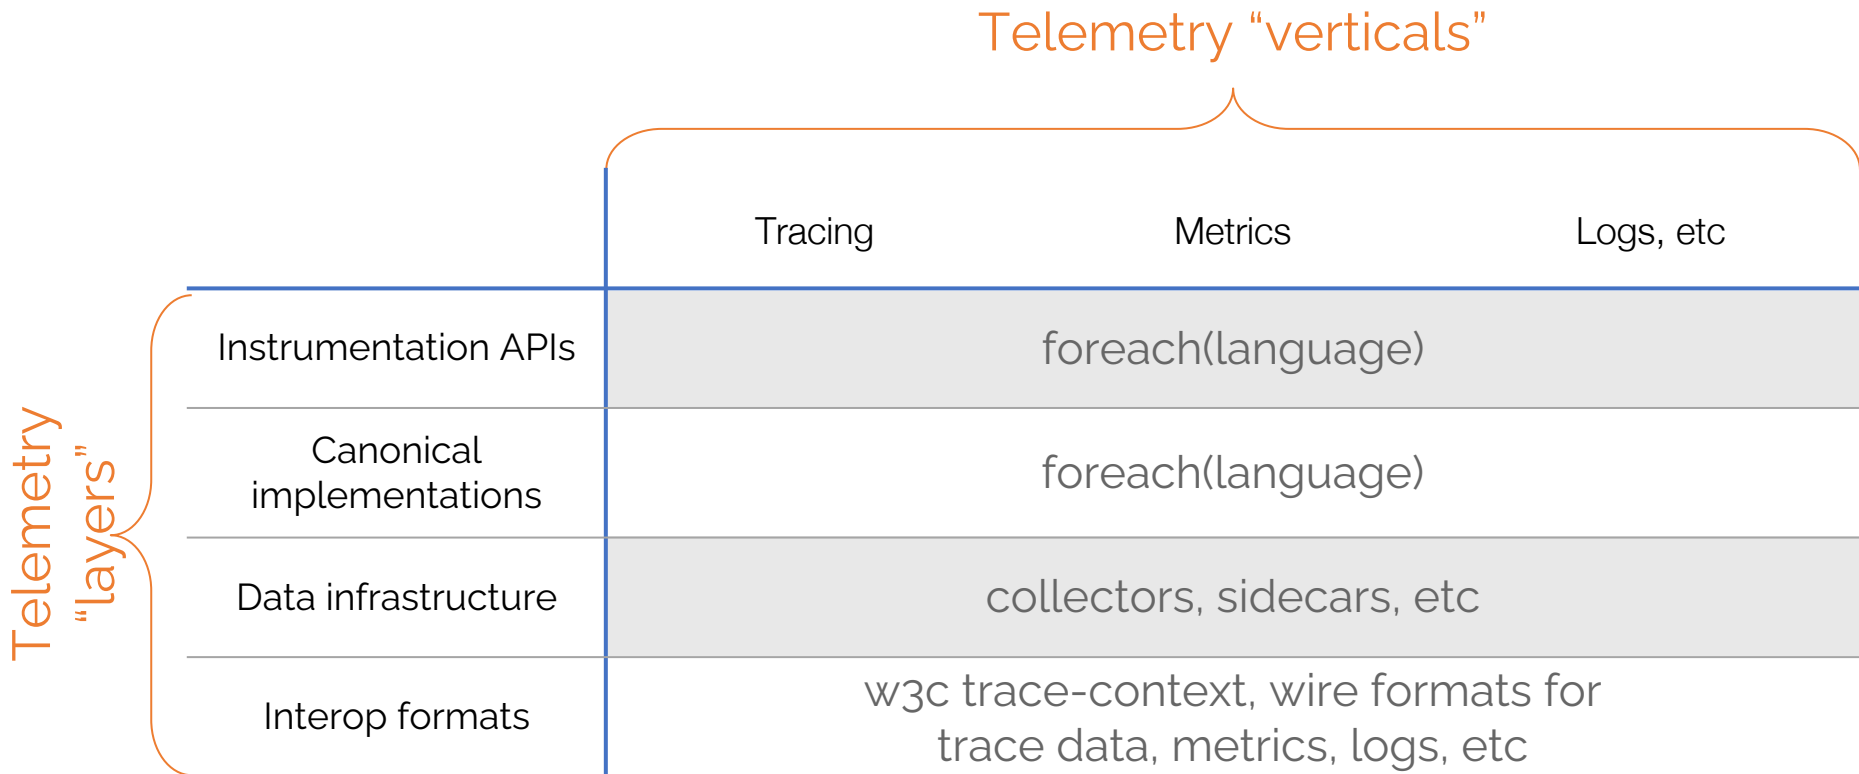
Mailchimp (PHP)
Postmates (Erlang)
Shopify (Ruby)

Other

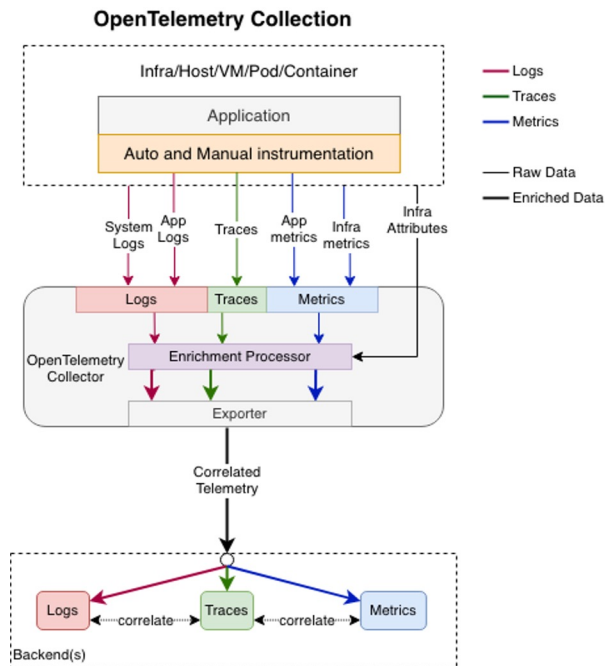



Jaeger > OtelCol
Fluent-bit <3 log SIG
Envoy roadmap
OpenMetrics roadmap
Spring roadmap

Cloud Native Telemetry

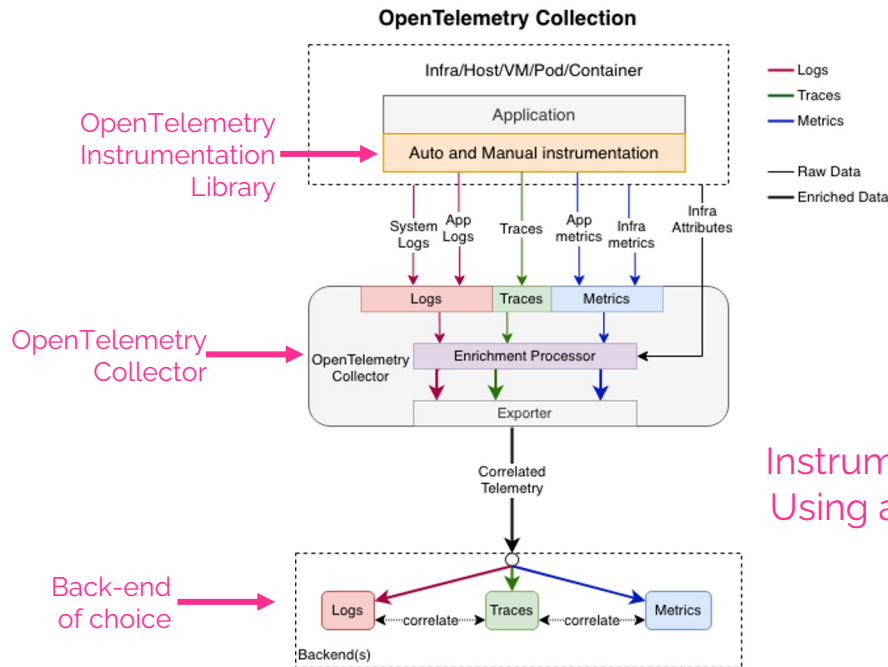


From Separate to Consolidate



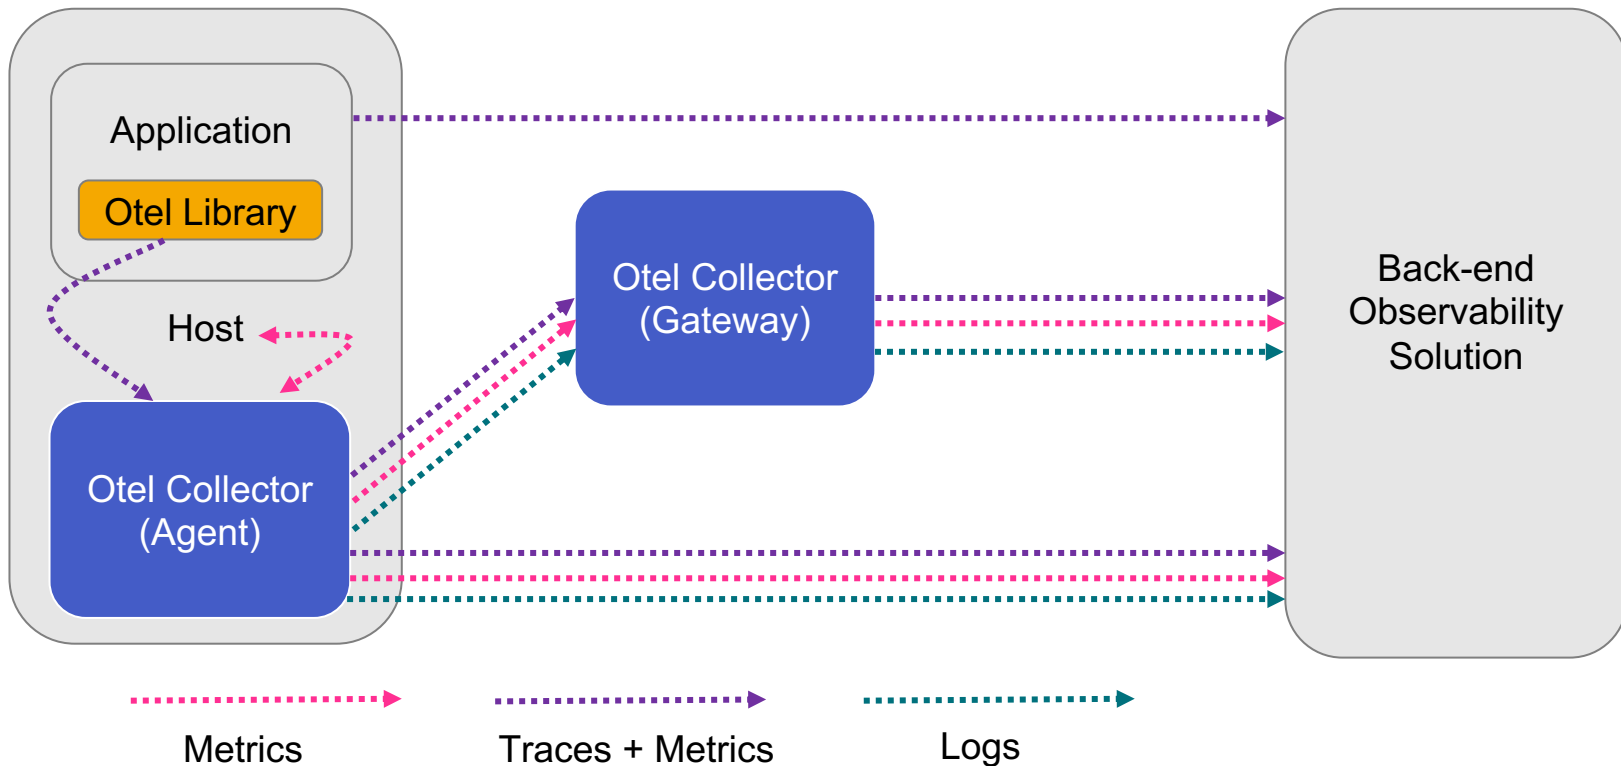
OpenTelemetry is NOT

- It's not an observability back-end like Jaeger or Prometheus.
- Instead it supports exporting data to a variety of open-source and commercial back-ends.

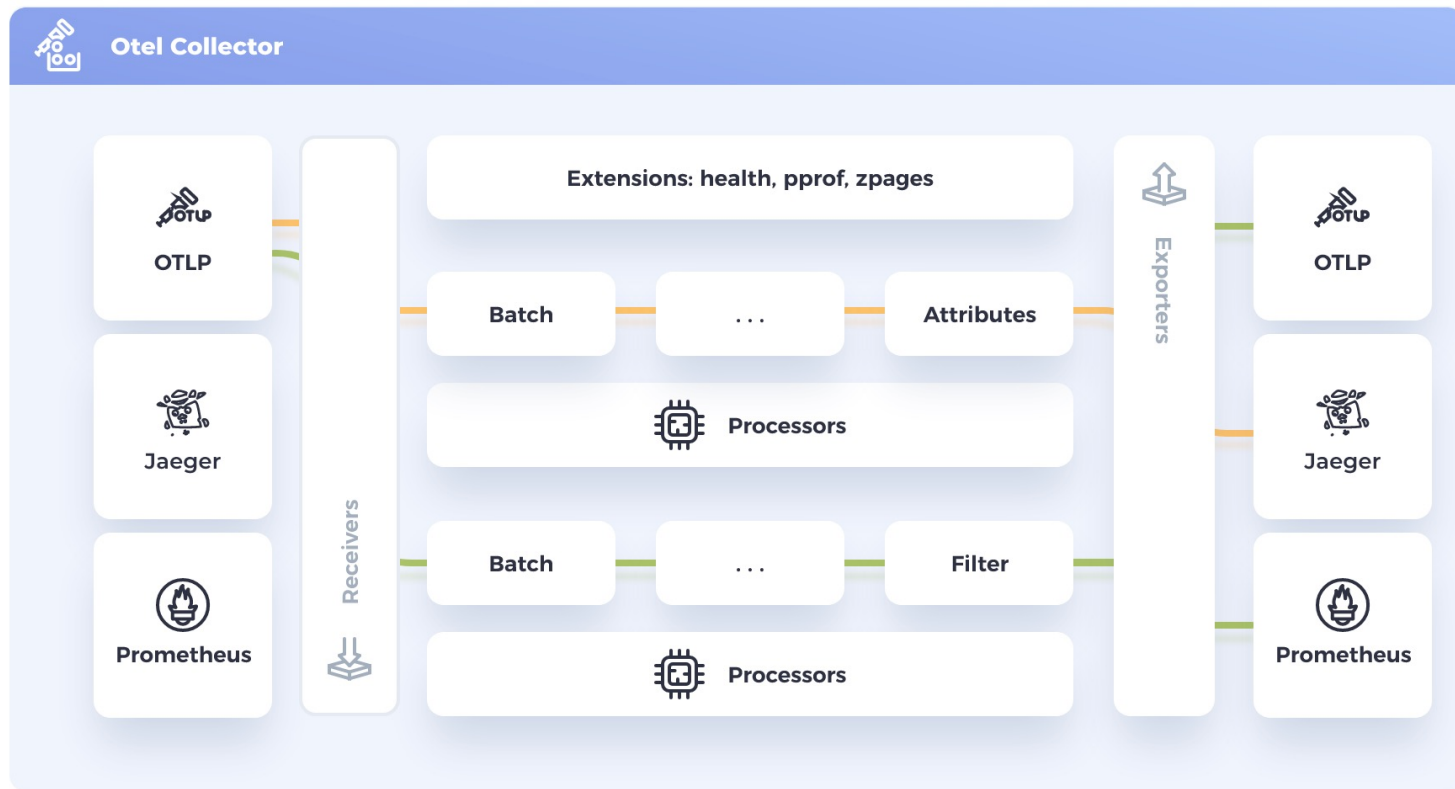


Before DEMO

Reference Architecture: OpenTelemetry



Reference Architecture: OpenTelemetry



Example of Core (Maintainers) Components

Traces

- Receivers/Exporters
 - OTLP
 - Jaeger
 - Zipkin
- Processors
 - Attributes
 - Batch
 - Queued Retry
 - Resource
 - Sampling
 - Span

Metrics

- Receivers
 - OTLP
 - Host (CPU, Disk, Memory, Network)
 - Prometheus
- Processors
 - *Coming soon...*
- Exporters
 - OTLP
 - Prometheus

Example of Contrib (Community) Components

Traces

- Processors
 - Kubernetes
- Exporters:
 - Alibaba
 - AWS X-ray
 - Azure Monitor
 - Elastic
 - Honeycomb
 - Kinesis
 - Lightstep
 - New Relic
 - Splunk
 - Stackdriver

Metrics

- Receivers
 - Carbon
 - Kubernetes
 - Redis
 - Splunk
 - Wavefront
- Processors
 - Metrics Transform
- Exporters
 - Carbon
 - Splunk
 - Stackdriver

Client Libraries: Java

Getting Started

Traces

1. Instantiate a tracer
2. Create spans
3. Enhance spans
4. Configure SDK

Metrics

1. Instantiate a meter
2. Create metrics
3. Enhance metrics
4. Configure observer

Getting Started: Traces (Manual)

```
# Instantiate tracer
Tracer tracer =
    OpenTelemetry.getTracer("instrumentation-library-name","semver:1.0.0");

# Create span
Span span = tracer.spanBuilder("my span").startSpan();
try (Scope scope = tracer.withSpan(span)) {
    // your use case
    # Enhance span
    span.setAttribute("version", "1.2");
} catch (Throwable t) {
    Status status = Status.UNKNOWN.withDescription("Change it to your error message");
    span.setStatus(status);
} finally {
    span.end(); // closing the scope does not end the span, this has to be done manually
}
```



Getting Started: Traces (Manual)

```
# Instantiate tracer
```

```
Tracer tracer =  
    OpenTelemetry.getTracer("instrumentation-library-name", "semver:1.0.0");
```

```
# Create span
```

```
Span span = tracer.spanBuilder("my span").startSpan();
```

```
try (Scope scope = tracer.withSpan(span)) {  
    // your use case  
    # Enhance span  
    span.setAttribute("version", "1.2");  
} catch (Throwable t) {  
    Status status = Status.UNKNOWN.withDescription("Change it to your error message");  
    span.setStatus(status);  
} finally {  
    span.end(); // closing the scope does not end the span, this has to be done manually  
}
```



Getting Started: Traces (Manual)

```
# Instantiate tracer
Tracer tracer =
    OpenTelemetry.getTracer("instrumentation-library-name","semver:1.0.0");

# Create span
Span span = tracer.spanBuilder("my span").startSpan();
try (Scope scope = tracer.withSpan(span)) {
    // your use case
    # Enhance span
    span.setAttribute("version", "1.2");
} catch (Throwable t) {
    Status status = Status.UNKNOWN.withDescription("Change it to your error message");
    span.setStatus(status);
} finally {
    span.end(); // closing the scope does not end the span, this has to be done manually
}
```

Getting Started: Traces (Manual)

```
# Instantiate tracer
Tracer tracer =
    OpenTelemetry.getTracer("instrumentation-library-name","semver:1.0.0");

# Create span
Span span = tracer.spanBuilder("my span").startSpan();
try (Scope scope = tracer.withSpan(span)) {
    // your use case
    # Enhance span
    span.setAttribute("version", "1.2");
} catch (Throwable t) {
    Status status = Status.UNKNOWN.withDescription("Change it to your error message");
    span.setStatus(status);
} finally {
    span.end(); // closing the scope does not end the span, this has to be done manually
}
```



Getting Started: Traces (Manual)

```
// Get the tracer
TracerSdkProvider tracerProvider = OpenTelemetrySdk.getTracerProvider();

// Configure the sampler to use
tracerProvider.updateActiveTraceConfig(
    TraceConfig alwaysOn = TraceConfig.getDefault().toBuilder().setSampler(
        Samplers.alwaysOn()
    ).build();
);

// Set to export the traces to via Jaeger
ManagedChannel jaegerChannel =
    ManagedChannelBuilder.forAddress([ip:String], [port:int]).usePlaintext().build();
JaegerGrpcSpanExporter jaegerExporter = JaegerGrpcSpanExporter.newBuilder()
    .setServiceName("example").setChannel(jaegerChannel).setDeadline(30000)
    .build();
tracerProvider.addSpanProcessor(
    BatchSpansProcessor.newBuilder(
        jaegerExporter
    ).build());
```



Getting Started: Traces (Manual)

```
// Get the tracer
TracerSdkProvider tracerProvider = OpenTelemetrySdk.getTracerProvider();

// Configure the sampler to use
tracerProvider.updateActiveTraceConfig(
    TraceConfig alwaysOn = TraceConfig.getDefault().toBuilder().setSampler(
        Samplers.alwaysOn()
    ).build();
);

// Set to export the traces to via Jaeger
ManagedChannel jaegerChannel =
    ManagedChannelBuilder.forAddress([ip:String], [port:int]).usePlaintext().build();
JaegerGrpcSpanExporter jaegerExporter = JaegerGrpcSpanExporter.newBuilder()
    .setServiceName("example").setChannel(jaegerChannel).setDeadline(30000)
    .build();
tracerProvider.addSpanProcessor(
    BatchSpansProcessor.newBuilder(
        jaegerExporter
    ).build());
```



Getting Started: Traces (Manual)

```
// Get the tracer
TracerSdkProvider tracerProvider = OpenTelemetrySdk.getTracerProvider();

// Configure the sampler to use
tracerProvider.updateActiveTraceConfig(
    TraceConfig.alwaysOn = TraceConfig.getDefault().toBuilder().setSampler(
        Samplers.alwaysOn()
    ).build();
);

// Set to export the traces to via Jaeger
ManagedChannel jaegerChannel =
    ManagedChannelBuilder.forAddress([ip:String], [port:int]).usePlaintext().build();
JaegerGrpcSpanExporter jaegerExporter = JaegerGrpcSpanExporter.newBuilder()
    .setServiceName("example").setChannel(jaegerChannel).setDeadline(30000)
    .build();
tracerProvider.addSpanProcessor(
    BatchSpansProcessor.newBuilder(
        jaegerExporter
    ).build());
```



Getting Started: Traces (Manual)

```
// Get the tracer
TracerSdkProvider tracerProvider = OpenTelemetrySdk.getTracerProvider();

// Configure the sampler to use
tracerProvider.updateActiveTraceConfig(
    TraceConfig alwaysOn = TraceConfig.getDefault().toBuilder().setSampler(
        Samplers.alwaysOn()
    ).build();
);

// Set to export the traces to via Jaeger
ManagedChannel jaegerChannel =
    ManagedChannelBuilder.forAddress([ip:String], [port:int]).usePlaintext().build();
JaegerGrpcSpanExporter jaegerExporter = JaegerGrpcSpanExporter.newBuilder()
    .setServiceName("example").setChannel(jaegerChannel).setDeadline(30000)
    .build();
tracerProvider.addSpanProcessor(
    BatchSpansProcessor.newBuilder(
        jaegerExporter
    ).build());
```



Getting Started: Traces (Manual)


```
// Get the tracer
TracerSdkProvider tracerProvider = OpenTelemetrySdk.getTracerProvider();

// Configure the sampler to use
tracerProvider.updateActiveTraceConfig(
    TraceConfig alwaysOn = TraceConfig.getDefault().toBuilder().setSampler(
        Samplers.alwaysOn()
    ).build();
);

// Set to export the traces to via Jaeger
ManagedChannel jaegerChannel =
    ManagedChannelBuilder.forAddress([ip:String], [port:int]).usePlaintext().build();
JaegerGrpcSpanExporter jaegerExporter = JaegerGrpcSpanExporter.newBuilder()
    .setServiceName("example").setChannel(jaegerChannel).setDeadline(30000)
    .build();
tracerProvider.addSpanProcessor(
    BatchSpansProcessor.newBuilder(
        jaegerExporter
    ).build());
```



Still with me?




There must be an
easier way...



WHAT IF I TOLD YOU

**YOU CAN INSTRUMENT
WITHOUT CODE CHANGES**



Getting Started: Traces (Automatic)

```
java -javaagent:path/to/opentelemetry-auto-<version>.jar \  
-Dota.exporter.jar=path/to/opentelemetry-auto-exporters-otlp-<version>.jar \  
-Dota.exporter.otlp.endpoint=localhost:55680 \  
-Dota.exporter.otlp.service.name=shopping \  
-jar myapp.jar
```



- Instruments known libraries with no code (only runtime) changes
- Adheres to semantic conventions
- Configurable via environment and/or runtime variables
- Can co-exist with manual instrumentation

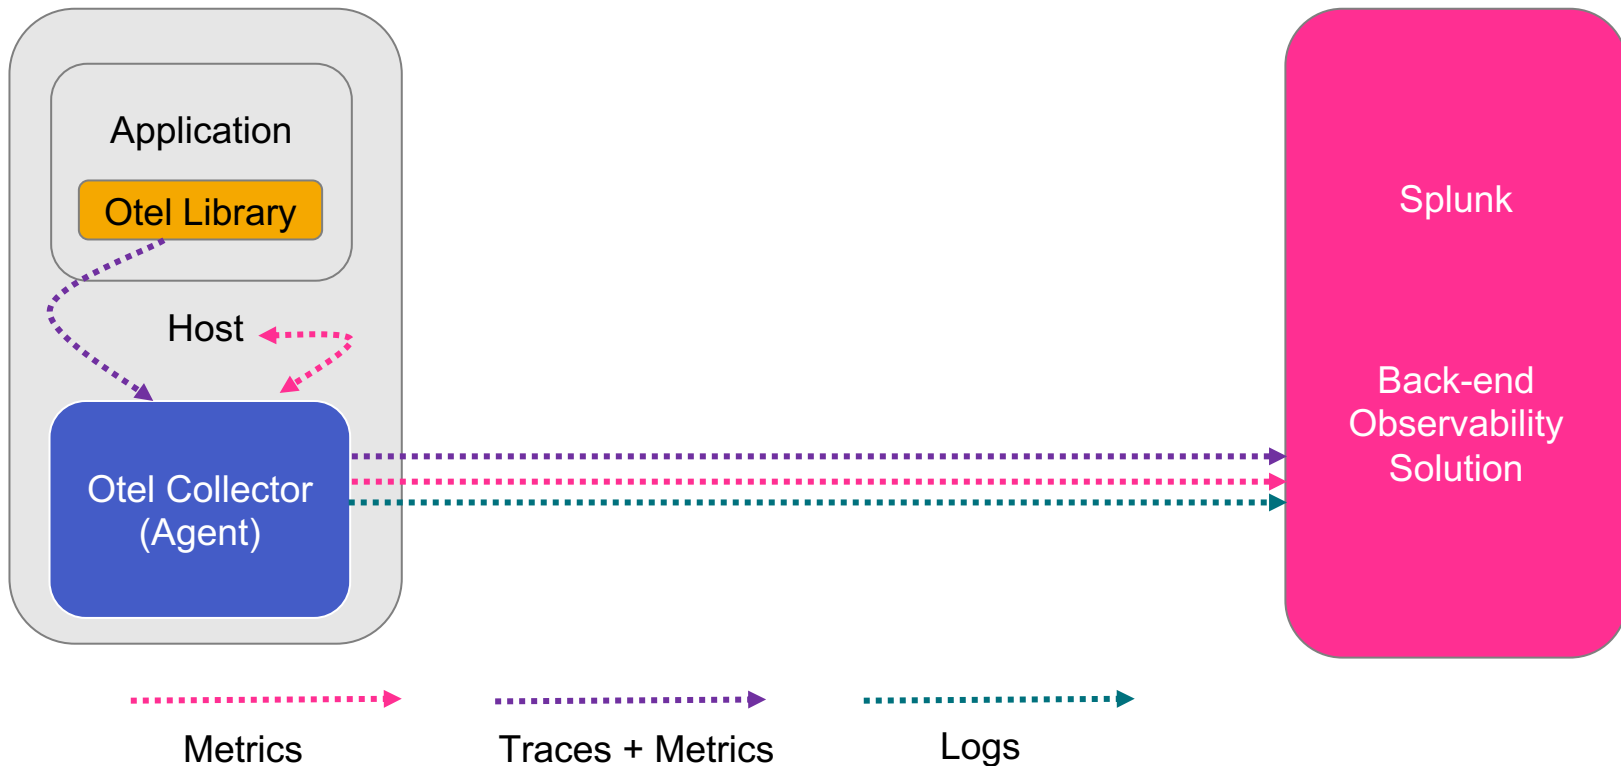
WARNING: Do not use two different auto-instrumentation solutions on the same service.

Akka HTTP 10.0+	Grizzly 2.0+	JSP 2.3+	Reactor 3.1+
Apache HttpClient 4.0+	gRPC 1.5+	Kafka 0.11+	Redis 1.8+
Apache HttpClient 2.0+	Hibernate 3.3+	Lettuce 4.0+	RMI Java 7+
AWS SDK 1.11.x and 2.2.0+	URLConnection Java 7+	Log4j 1.1+	RxJava 1.0+
Cassandra Driver 3.0+ (not 4.x yet)	Hystrix 1.4+	Logback 1.0+	Servlet 2.3+
Couchbase Client 2.0+ (not 3.x yet)	Java.util.logging Java 7+	MongoDB Drivers 3.3+	Spark Web Framework 2.3+
Dropwizard Views 0.7+	JAX-RS 0.5+	Netty 3.8+	Spring Data 1.8+
Elasticsearch API 2.0+ (not 7.x yet)	JAX-RS Client 2.0+	OkHttp 3.0+	Spring Scheduling 3.1+
Elasticsearch REST Client 5.0+	JDBC Java 7+	Play 2.3+ (not 2.8.x yet)	Spring Servlet MVC 3.1+
Finatra 2.9+	Jedis 1.4+	Play WS 1.0+	Spring Webflux 5.0+
Geode Client 1.4+	Jetty 8.0+	RabbitMQ Client 2.7+	Spymemcached 2.12+
Google HTTP Client 1.19+	JMS 1.1+	Ratpack 1.5+	Twilio 6.6+

OpenTelemetry Java auto-instrumentation library support

DEMO

Demo Architecture: OpenTelemetry



Links

- Specification
 - <https://github.com/open-telemetry/opentelemetry-specification>
- OpenTelemetry Collector
 - <https://opentelemetry.io/docs/collector/about/>
 - <https://opentelemetry.io/docs/collector/configuration/>
- Java client library
 - <https://github.com/open-telemetry/opentelemetry-java/blob/master/QUICKSTART.md>
 - <https://github.com/open-telemetry/opentelemetry-auto-instr-java>
- Other
 - <https://opentelemetry.io/docs/workshop/resources/>
 - <https://devstats.cncf.io/>
 - <https://medium.com/jaegertracing/jaeger-embraces-opentelemetry-collector-90a545cbc24>
 - <https://github.com/spring-petclinic/spring-petclinic-microservices>

Thank You!

Additional resources:

- <https://opentelemetry.io>
- <https://github.com/open-telemetry/community>
- <https://www.cncf.io/webinars/how-opentelemetry-is-eating-the-world/>